Computational Structures in Data Science

Lecture #14: Performance and Distributed Computing

UC Berkeley EECS
Adj. Assistant Prof.
Dr. Gerald Friedland

MIT Technology Review

Intelligent Machines

With this tool, AI could identify new malware as readily as it recognizes cats

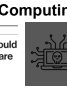

https://www.technologyreview.com/s/610881/with-this-tool-ai-could-identify-malware-as-readily-as-it-recognizes-cats/

April 20, 2018 http://inst.eecs.berkeley.edu/~cs88

Administrivia

- This is the last "required" lecture. Next week: Information and bits, Summary.
- The week after: Q&A for finals.
- Today: HKN review!
  Please do the survey and give us good grades! ☺
- Thank you:
  - TAs!
  - Lab Assistants!
  - UC Berkeley Staff!

Computational Concepts Toolbox

- Data type: values, literals, operations,
- Expressions, Call expression
- Variables
- Assignment Statement
- Sequences: tuple, list
- Dictionaries
- Data structures
- Tuple assignment
- Function Definition Statement
- Conditional Statement
- Iteration: list comp, for, while
- Lambda function expr.
- Higher Order Functions
  - as Values, Args, Results
- Higher order function patterns
  - Map, Filter, Reduce
  - Function factories
- Recursion
  - Linear, Tail, Tree
- Abstract Data Types
- Mutation
- Iterators and Generators
- Object Oriented Programming, Classes
- Exceptions
- Declarative Programming
- Distributed Computing

Recap: Complexity

- Example: Matrix Multiply
  - How many Multiplies? Adds? Ops? How much time ?
  - As a function of n ?

```
for i in 0 to n-1:
    for j in 0 to n-1:
        C[i][j] = 0
        for k in 0 to n-1:
            C[i][j] = C[i][j] + A[i][k]*B[k][j]
```

We say it is $O(n^3)$ "big-O of n^3" as an *asymptotic upper bound*

$time(n) < c \cdot n^3$, for some suitably large constant *c* for any instance of the inputs of size n.

A more subtle complexity example

- What is the "complexity" of finding the average number of factors of numbers up to n?

```
def factors(n):
  return [x for x in range(2, max(n, ceil(sqrt(n))))
                                      if n % x == 0]

def ave_factor(n):
    all_factors = map(factors, range(n))
    all_lens = map(len, all_factors)
    return sum(all_lens)/n
```

(annotations: $n^{1/2}$, n)

```
from timeit import default_timer as timer

def timeit(fun):
    """ Rtn timer for fun(i) in secs. """
    def timer_fun(i):
        start = timer()
        fun(i)
        end = timer()
        return (end-start)
    return timer_fun
```


How long does factors take?

```
In [9]: tbl = Table().with_column('n', np.arange(0,1000, 1))
        tbl['factors'] = tbl.apply(factors, 'n')
        tbl['n_factors'] = tbl.apply(len, 'factors')
        tbl['secs'] = tbl.apply(timeit(factors), 'n')
        tbl
```

Out[9]:

n	factors	n_factors	secs
0	[]	0	9.76503e-06
1	[]	0	2.40898e-06
2	[]	0	1.34797e-06
3	[]	0	3.49898e-06
4	[2]	1	2.74903e-06
5	[]	0	2.43704e-06
6	[2, 3]	2	3.019e-06
7	[]	0	2.78e-06
8	[2, 4]	2	3.28396e-06
9	[3]	1	3.74601e-06

... (990 rows omitted)

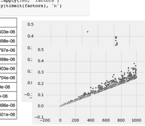

Big Data, Big Problems

- Performance terminology
  - the FLOP: FLoating point OPeration
  - "flops" = # FLOP/second is the standard metric for computing power
- Example: Global Climate Modeling
  - Divide the world into a grid (e.g. 10 km spacing)
  - Solve fluid dynamics equations for each point & minute
    - » Requires about 100 Flops per grid point per minute
  - Weather Prediction (7 days in 24 hours):
    - » 56 Gflops
  - Climate Prediction (50 years in 30 days):
    - » 4.8 Tflops
- Perspective
  - Intel Core i7 980 XE Desktop Processor
    - » ~100 Gflops
    - » Climate Prediction would take ~5 years

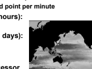

www.epm.ornl.gov/chammp/chammp.html

What Can We Do? Use Many CPUs!

- **Supercomputing** – like those listed in top500.org
  - Multiple processors "all in one box / room" from one vendor that often communicate through shared memory
  - This is often where you find exotic architectures
- **Distributed computing**
  - Many separate computers (each with independent CPU, RAM, HD, NIC) that communicate through a network
    - » Grids (heterogenous computers across Internet)
    - » Clusters (mostly homogeneous computers all in one room)
      - Google uses commodity computers to exploit "knee in curve" price/performance sweet spot
  - It's about being able to solve "big" problems, not "small" problems faster
    - » These problems can be data (mostly) or CPU intensive

Recap: Filter, Map, Reduce

- Functions as Data
- Higher-Order Functions
- Useful HOFs (you can build your own!)
  - map function over list
    - » Report a new list, every element e of list becoming function(e)
  - filter items such that predicate from list
    - » Create a new list, keeping only elements e of list if predicate(e)
  - reduce with function over list
    - » Combine all the elements of list with function(e)

- Example:
  filter → map → reduce

en.wikipedia.org/wiki/MapReduce

Google's MapReduce Simplified

- Filter: Chunk data and send to different CPUs.
- Map: Apply function to data chunks on different CPUs.
- Reduce: Combine results from different CPUs.
  - Reducer should be associative and commutative
- Imagine 10,000 machines ready to help you compute anything you could cast as a MapReduce problem!
  - This is the abstraction Google is famous for authoring
  - The system takes care of load balancing, dead machines, etc.

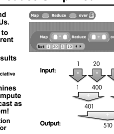

MapReduce: Advantages/Disadvantages

- Now it's easy to program for many CPUs
  - Communication management effectively gone
  - Fault tolerance, monitoring
    - » machine failures, suddenly-slow machines, etc are handled
  - Can be much easier to design and program!
  - Can cascade several (many?) MapReduce tasks
- But… it might restrict solvable problems
  - Might be hard to express problem in MapReduce
  - Data parallelism is key
    - » Need to be able to break up a problem by data chunks

Apache Spark (from Berkeley)

- Data processing system that provides a simple interface to analytics on large data
- A Resilient Distributed Dataset (RDD) is a collection of values or key-value pairs
- Support the operations you are familiar with
  - Data-Parallel: map, filter, reduce
  - Database: join, union, intersect
  - OS: sort, distinct, count
- All of can be performed on RDDs that are partitioned across machines

Spark Execution Model

Processing is defined centrally and executed remotely

- A RDD is distributed over workers
- A driver program defines transformations and actions on RDDs
- A cluster manager assigns task to workers
- Workers perform computation, store data, & communicate with each other
- Final results communicate back to driver

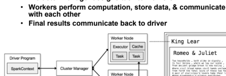

Distributed Computing Challenges

- Communication is fundamental difficulty
  - Distributing data, updating shared resource, communicating results, handling failures
  - Machines have separate memories, so need network
  - Introduces inefficiencies: overhead, waiting, etc.
- Need to parallelize algorithms, data structures
  - Must look at problems from parallel standpoint
  - Best for problems whose compute times >> overhead

Speedup Issues: Amdahl's Law

- Applications can almost never be completely parallelized; some serial code remains

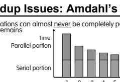

- s is serial fraction of program, P is # of cores (was processors)
- **Amdahl's law:**

Speedup(P) = Time(1) / Time(P)

$\leq 1 / (s + [(1-s) / P)]$, and as $P \to \infty$

$\leq 1 / s$

Amdahl's Law: Conclusion

- **Computer Science View**: Even if the parallel portion of your application speeds up perfectly, your performance will be limited by the sequential portion. ☹
- **Data Science View:** Often, as the data gets large, the work that can be parallelized grows faster than the size of the data. ☺

Fundamental Change in Perspective!

Summary: Data science

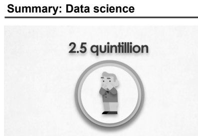

https://www.youtube.com/watch?v=TzxmjbL-i4Y

Summary: Distributed Computing

- Parallelization can help speed up
- Bottleneck: dependencies in data, algorithm
- Requires rethinking the program
- Good luck with the project!
- See you next week for that final (fun) lecture!