




UC Berkeley EECS  
Adj. Assistant Prof.  
Dr. Gerald Friedland

## MIT Technology Review

# Lecture #14: Performance and Distributed Computing

Intelligent Machines

With this tool, AI could identify new malware as readily as it recognizes cats



<https://www.technologyreview.com/s/610881/with-this-tool-ai-could-identify-malware-as-readily-as-it-recognizes-cats/>

April 20, 2018

<http://inst.eecs.berkeley.edu/~cs88>

## Administrivia



- This is the last “required” lecture. Next week: Information and bits, Summary.
- The week after: Q&A for finals.
- Today: HKN review!  
Please do the survey and give us good grades! ☺
- Thank you:
  - TAs!
  - Lab Assistants!
  - UC Berkeley Staff!

20/04/18

UCB CS88 Sp18 L14

2

## Computational Concepts Toolbox



- Data type: values, literals, operations,
- Expressions, Call expression
- Variables
- Assignment Statement
- Sequences: tuple, list
- Dictionaries
- Data structures
- Tuple assignment
- Function Definition Statement
- Conditional Statement
- Iteration: list comp, for, while
- Lambda function expr.
- Higher Order Functions
  - as Values, Args, Results
- Higher order function patterns
  - Map, Filter, Reduce
  - Function factories
- Recursion
  - Linear, Tail, Tree
- Abstract Data Types
- Mutation
- Iterators and Generators
- Object Oriented Programming, Classes
- Exceptions
- Declarative Programming
- **Distributed Computing**

## Recap: Complexity



- **Example: Matrix Multiply**
  - How many Multiplies? Adds? Ops? How much time ?
  - As a function of  $n$  ?

```
for i in 0 to n-1:
  for j in 0 to n-1:
    C[i][j] = 0
    for k in 0 to n-1:
      C[i][j] = C[i][j] + A[i][k]*B[k][j]
```

We say it is  $O(n^3)$  “big-O of  $n^3$ ” as an *asymptotic upper bound*

$\text{time}(n) < c \cdot n^3$ , for some suitably large constant  $c$  for any instance of the inputs of size  $n$ .



20/04/18

UCB CS88 Sp18 L14

3

20/04/18

UCB CS88 Sp18 L14

4

## A more subtle complexity example

- What is the “complexity” of finding the average number of factors of numbers up to  $n$ ?

```
def factors(n):  
    return [x for x in range(2, max(n, ceil(sqrt(n))))  
            if n % x == 0]  
  
def ave_factor(n):  
    all_factors = map(factors, range(n))  
    all_lens = map(len, all_factors)  
    return sum(all_lens)/n
```

```
from timeit import default_timer as timer  
  
def timeit(fun):  
    """ Rtn timer for fun(i) in secs. """  
    def timer_fun(i):  
        start = timer()  
        fun(i)  
        end = timer()  
        return (end-start)  
    return timer_fun
```

20/04/18

UCB CS88 Sp18 L14

5


## How long does factors take?

```
In [9]: tbl = Table().with_column('n', np.arange(0,1000, 1))  
tbl['factors'] = tbl.apply(factors, 'n')  
tbl['n_factors'] = tbl.apply(len, 'factors')  
tbl['secs'] = tbl.apply(timeit(factors), 'n')  
tbl
```

```
Out[9]:
```

n	factors	n_factors	secs
0	[]	0	9.76503e-06
1	[]	0	2.40898e-06
2	[]	0	1.34797e-06
3	[]	0	3.49898e-06
4	[2]	1	2.74903e-06
5	[]	0	2.43704e-06
6	[2, 3]	2	3.019e-06
7	[]	0	2.78e-06
8	[2, 4]	2	3.28396e-06
9	[3]	1	3.74601e-06

... (990 rows omitted)




20/04/18

UCB CS88 Sp18 L14

6

## Big Data, Big Problems

- Performance terminology**
  - the FLOP: Floating point Operation
  - “flops” = # FLOP/second is the standard metric for computing power
- Example: Global Climate Modeling**
  - Divide the world into a grid (e.g. 10 km spacing)
  - Solve fluid dynamics equations for each point & minute
    - Requires about 100 Flops per grid point per minute
  - Weather Prediction (7 days in 24 hours):
    - 56 Gflops
  - Climate Prediction (50 years in 30 days):
    - 4.8 Tflops
- Perspective**
  - Intel Core i7 980 XE Desktop Processor
    - ~100 Gflops
    - Climate Prediction would take ~5 years



www.epm.ornl.gov/chammp/chammp.html

## What Can We Do? Use Many CPUs!

- Supercomputing** – like those listed in top500.org
  - Multiple processors “all in one box / room” from one vendor that often communicate through shared memory
  - This is often where you find exotic architectures
- Distributed computing**
  - Many separate computers (each with independent CPU, RAM, HD, NIC) that communicate through a network
    - Grids (heterogenous computers across Internet)
    - Clusters (mostly homogeneous computers all in one room)
      - Google uses commodity computers to exploit “knee in curve” price/performance sweet spot
  - It’s about being able to solve “big” problems, not “small” problems faster
    - These problems can be data (mostly) or CPU intensive


20/04/18

UCB CS88 Sp18 L14

8

## Recap: Filter, Map, Reduce

- Functions as Data
- Higher-Order Functions
- Useful HOFs (you can build your own!)
  - **map function over list**
    - » Report a new list, every element  $e$  of list becoming  $\text{function}(e)$
  - **filter items such that predicate from list**
    - » Create a new list, keeping only elements  $e$  of list if  $\text{predicate}(e)$
  - **reduce with function over list**
    - » Combine all the elements of list with  $\text{function}(e)$
- Example:
  - filter  $\rightarrow$  map  $\rightarrow$  reduce



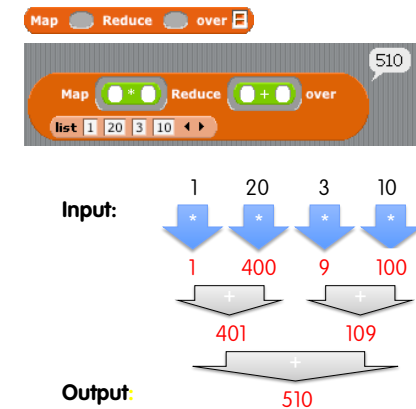
20/04/18

UCB CS88 Sp18 L14

9

## Google's MapReduce Simplified

- Filter: Chunk data and send to different CPUs.
- Map: Apply function to data chunks on different CPUs.
- Reduce: Combine results from different CPUs.
  - Reducer should be associative and commutative
- Imagine 10,000 machines ready to help you compute anything you could cast as a MapReduce problem!
  - This is the abstraction Google is famous for authoring
  - The system takes care of load balancing, dead machines, etc.



20/04/18

UCB CS88 Sp18 L14

10

## MapReduce: Advantages/Disadvantages

- Now it's easy to program for many CPUs
  - Communication management effectively gone
  - Fault tolerance, monitoring
    - » machine failures, suddenly-slow machines, etc are handled
  - Can be much easier to design and program!
  - Can cascade several (many?) MapReduce tasks
- But... it might restrict solvable problems
  - Might be hard to express problem in MapReduce
  - Data parallelism is key
    - » Need to be able to break up a problem by data chunks

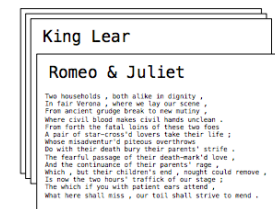
20/04/18

UCB CS88 Sp18 L14

11

## Apache Spark (from Berkeley)

- Data processing system that provides a simple interface to analytics on large data
- A Resilient Distributed Dataset (RDD) is a collection of values or key-value pairs
- Support the operations you are familiar with
  - Data-Parallel: map, filter, reduce
  - Database: join, union, intersect
  - OS: sort, distinct, count
- All of can be performed on RDDs that are partitioned across machines



20/04/18


UCB CS88 Sp18 L14

12

## Spark Execution Model

Processing is defined centrally and executed remotely

- A RDD is distributed over workers
- A driver program defines transformations and actions on RDDs
- A cluster manager assigns task to workers
- Workers perform computation, store data, & communicate with each other
- Final results communicate back to driver



20/04/18

UCB CS88 Sp18 L14

13

## Distributed Computing Challenges

- Communication is fundamental difficulty
  - Distributing data, updating shared resource, communicating results, handling failures
  - Machines have separate memories, so need network
  - Introduces inefficiencies: overhead, waiting, etc.
- Need to parallelize algorithms, data structures
  - Must look at problems from parallel standpoint
  - Best for problems whose compute times  $\gg$  overhead


20/04/18

UCB CS88 Sp18 L14

14

## Speedup Issues: Amdahl's Law

- Applications can almost never be completely parallelized; some serial code remains



- $s$  is serial fraction of program,  $P$  is # of cores (was processors)
- **Amdahl's law:**

$$\text{Speedup}(P) = \text{Time}(1) / \text{Time}(P)$$

$$\leq 1 / (s + [(1-s) / P]), \text{ and as } P \rightarrow \infty$$

$$\leq 1 / s$$

## Amdahl's Law: Conclusion

- **Computer Science View:** Even if the parallel portion of your application speeds up perfectly, your performance will be limited by the sequential portion. ☹️

- **Data Science View:** Often, as the data gets large, the work that can be parallelized grows faster than the size of the data. 😊

**Fundamental Change in Perspective!**

20/04/18

UCB CS88 Sp18 L14

15

20/04/18

UCB CS88 Sp18 L14

16

## Summary: Data science



<https://www.youtube.com/watch?v=TzxmjbL-i4Y>

## Summary: Distributed Computing



- Parallelization can help speed up
  - Bottleneck: dependencies in data, algorithm
  - Requires rethinking the program
- 
- Good luck with the project!
  - See you next week for that final (fun) lecture!